

#Structuring the notebooks based on the collected and parsed metadata

```
ask_structured_schema = {
    "name": "ask_structured_schema",
    "description": (
        "***IMPORTANT**": Your *entire* response must be valid JSON matching this
        schema—**no** single-quotes, no Python `None`, no trailing commas, no code fences,\n"
        "From the competition metadata, dataset metadata, and a raw Jupyter notebook text, "
        "extract exactly these fields as JSON (no extra keys, no prose, no markdown):\n"
        " - competition_problem_type: one of ['classification','regression']\n"
        " - competition_problem_subtype: single, concise, lowercase-and-hyphenated phrase
        (e.g. “binary-classification”, “multiclass-classification”, “multi-label-classification”,
        “time-series-forecasting”, “continuous-regression”, “ordinal-regression”, etc. or any other that
        fits.)\n"
        " - competition_problem_description: dense, short, factual description of the problem, what
        needs to be found, no repetitive words (omit dataset-meta here)\n"
        " - dataset_metadata: plain-English dataset_metadata in plain English as a single
        coherent paragraph, removing any non-human symbols (no bullets or symbols)\n"
        " - competition_dataset_type: one of
        ['Tabular','Time-series','Text','Image','Audio','Video','Geospatial','Graph','Multimodal']\n"
        " - preprocessing_steps: array of strings, each describing one transformation (e.g.
        'median-impute missing values')\n"
        " - notebook_model_layers_code: literal code snippet that builds(e.g model.fit) each
        layer(e.g Dense, Conv, etc..) and compiles the model(e.g model.compile) \n"
        " - used_technique: either 'DL' or 'ML'\n"
        " - library: string naming the main library used (exactly one 'Tensorflow', 'Pytorch')\n"
        " - target_column: array of all column names in the dataset that must be predicted \n"
        "Emit ONLY those keys."
    ),
    "parameters": {
        "type": "object",
        "properties": {
            "competition_problem_type": {
                "type": "string",
                "enum": ["classification", "regression"],
                "description": "Pick exactly one."
            },
            "competition_problem_subtype": {
                "type": "string",
                "enum": [
                    "binary-classification",
                    "multiclass-classification",
                    "multi-label-classification",
                    "time-series-forecasting",
```

```

        "continuous-regression",
        "quantile-regression",
        "multi-output-regression",
        "ordinal-regression",
        "missing-value-imputation"
    ],
    "description": "Pick exactly one."
},
"competition_problem_description": {
    "type": "string",
    "description": "Dense, factual description of what needs to be predicted."
},
"dataset_metadata": {
    "type": "string",
    "description": "Plain-English paragraph describing the dataset."
},
"competition_dataset_type": {
    "type": "string",
    "enum":
["Tabular", "Time-series", "Text", "Image", "Audio", "Video", "Geospatial", "Graph", "Multimodal"],
    "description": "Choose one primary modality."
},
"preprocessing_steps": {
    "type": "array",
    "items": {"type": "string"},
    "description": "List every transformation (scaling, normalization, one-hot encoding,
etc.) in plain English."
},
"notebook_model_layers_code": {
    "type": "string",
    "description": (
        "include the literal code lines of model compile, model fit, and that build each layer
(e.g. `Dense(128, activation='relu', ...)`)\\n"
        "The line(s) that create or instantiate the model (Sequential, Functional, subclass,
torch.nn.Module, etc.).\\n"
        "All layer-construction calls (Dense, Conv2D, custom layers, etc.) or layer
definitions in a subclass.\\n"
        "The call that compiles or configures training (e.g. `compile()`,
`configure_optimizers()`, etc.).\\n"
        "The call that launches training (e.g. `fit()`, `trainer.fit()`, `train()`, etc.).\\n"
        "Do not include unrelated code, helper wrappers, or omit any of these steps."
    )
},
"used_technique": {

```

```
    "type": "string",
    "enum": ["DL","ML"],
    "description": "Either 'DL' or 'ML'."
  },
  "library": {
    "type": "string",
    "description": "Name of the main library used."
  },
  "target_column": {
    "type": "array",
    "items": {"type": "string"},
    "description": "List of all column names in train.csv to predict."
  }
},
"required": [
  "competition_problem_type",
  "competition_problem_subtype",
  "competition_problem_description",
  "dataset_metadata",
  "competition_dataset_type",
  "preprocessing_steps",
  "notebook_model_layers_code",
  "used_technique",
  "library",
  "target_column"
]
}
```

#Building a Keras model using the metadata

```
tools = [
    {
        "name": "generate_keras_schema",
        "type": "function",
        "description": (
            """Generate and save a runnable deep learning model using Keras in Python code
            wrapped in <Code>...</Code> in a single `notebook_code` JSON field:\n"
            "The generated code must implement:\n"
            "1. **Reproducibility**: set seeds for Python, NumPy, scikit-learn, and TensorFlow (or
            PyTorch).\n"
            "2. **Imports**:\n" \n"
            " - `pandas`, `numpy`\n" \n"
            " - `sklearn.model_selection.train_test_split`\n" \n"
            " - `sklearn.impute.SimpleImputer`\n" \n"
            " - `sklearn.compose.ColumnTransformer`\n" \n"
            " - `sklearn.preprocessing.StandardScaler`, `OneHotEncoder`, `LabelEncoder` ←
            **added here**\n" \n"
            " - `sklearn.pipeline.Pipeline`\n" \n"
            " - `tensorflow` (or `torch`)\n" \n"
            " - `tensorflow.keras.callbacks.EarlyStopping, ModelCheckpoint`\n" \n"
            " - `from tensorflow.keras.metrics import SparseTopKCategoricalAccuracy`\n" \n"
            " - `json`, `time`\n" \n"
            " When using OneHotEncoding, use sparse_output=False instead of sparse\n"
            "3. Data Loading, Split & Target Encoding:\n"
            "Read each file in files_list into train_dfs\n"
            "If any filename ends with 'test.csv', load it into df_test, else df_test=None\n"
            "Infer id_col & target_columns from submission_example header\n"
            "df = pd.concat(train_dfs, ignore_index=True)\n"
            "# Target encoding immediately after df is final:\n"
            "col = target_columns[0]\n"
            "if competition_problem_subtype in ['binary-classification']:\n"
            "    from sklearn.preprocessing import LabelEncoder\n"
            "    le=LabelEncoder().fit(df[col].astype(str))\n"
            "    y_enc=le.transform(df[col].astype(str)).astype(int)\n"
            "    classes_=le.classes_\n"
            "elif competition_problem_subtype in ['multiclass-classification', 'multiclass
            classification', 'ordinal-regression']:\n"
            "    from sklearn.preprocessing import LabelEncoder\n"
            "    le=LabelEncoder().fit(df[col].astype(str))\n"
            "    y_enc=le.transform(df[col].astype(str))\n"
            "    classes_=le.classes_
```

```

"elif competition_problem_subtype in ['multi-label-classification']:\n"
"    from sklearn.preprocessing import MultiLabelBinarizer\n"
"    mlb=MultiLabelBinarizer()\n"
"    y_enc=mlb.fit_transform(df[target_columns])\n"
"    classes_=mlb.classes_\n"
"elif competition_problem_subtype in
['continuous-regression','quantile-regression','multi-output
regression','missing-value-imputation']:\n"
"    y_values = df[target_columns].astype(float).values\n"
"    y_enc = np.log1p(y_values) if np.all(y_values >= 0) else y_values\n"
"elif competition_problem_subtype in
['time-series-forecasting','multivariate-time-series-forecasting']:\n"
"    y_enc=df[target_columns].values\n"
"else:\n"
"    y_enc=df[target_columns].values\n"
"X=df.drop(columns=target_columns+[id_col],errors='ignore')\n"
"# now either use provided df_test or split off 20% for test:\n"
"if df_test is None:\n"
"    X_train,X_val,y_train,y_val=train_test_split(\n"
"        X,y_enc,\n"
"        test_size=0.2,\n"
"        stratify=y_enc if competition_problem_subtype in
['binary-classification','multiclass-classification','multiclass classification'] else None,\n"
"        random_state=42)\n"
"    train_ids=X_train[id_col]\n"
"    test_ids =X_val[id_col]\n"
"else:\n"
"    X_train=X\n"
"    y_train=y_enc\n"
"    train_ids=df[id_col]\n"
"    test_ids =df_test[id_col]\n"
"    X_val =df_test.drop(columns=target_columns+[id_col],errors='ignore')\n"
"    y_val = None # explicitly set\n"
"\n"
"4. Feature Engineering:\n"
"    Automatically drop columns with all missing values\n"
"    Identify categorical columns and remove those with extremely high cardinality (eg
>50 unique)\n"
"    Optionally apply any additional simple transformations you deem useful\n"
"5. **Preprocessing Pipeline**:\n"
"    - Auto-detect numeric vs. categorical via `df.select_dtypes`.\n"
"    - Build a `ColumnTransformer` with median-imputed & scaled numerics, and
most-frequent-imputed & OHE categoricals (cap at 50 cats).\n"
"    - Fit on train → transform train/val/test.\n"

```

```

"6. Model Architecture:"
"- Build at least two hidden layers with BatchNormalization and Dropout after each\n"
"- Set output units = number of target_columns for multilabel/multiclass, else 1\n"
"- Choose depth & width by data shape: shallow/narrow for small datasets,
deeper/wider for large datasets, scale units  $\approx \min(\text{features} \times 2, 1024)$ \n"
"- Leverage provided `examples` but adjust architecture based on dataset size,
feature count, and target count\n"
"- Architectural Guidelines:\n"
"  - Choose by data size:\n"
"    • If `n_samples < 10000` or `n_features < 100`:\n"
"      - Build two Dense layers of sizes:\n"
"        [min(n_features*2, 128), min(n_features, 64)]\n"
"      - No BatchNormalization; Dropout  $\leq 0.3$ \n"
"    • Else:\n"
"      - Build 2–4 Dense layers of sizes:\n"
"        [min(n_features*i, 1024) for i in (2, 1, 0.5, 0.25)] (drop any <16 units)\n"
"      - After each: BatchNormalization() + Dropout( $\leq 0.4$ )\n"
"\n"
For all hidden Dense layers (except the final output), use ReLU activation\n"
"  - Task subtype  $\rightarrow$  head, loss, batch & metrics:\n"
"    (Note: activation applies only to the final/output layer)\n"
"    * binary-classification:\n"
"      - activation=sigmoid, loss=binary_crossentropy\n"
"      - batch_size=64–256, metrics=['accuracy', tf.keras.metrics.AUC(),
tfa.metrics.MatthewsCorrelationCoefficient()]\n"
"    * multiclass-classification (MAP@N):\n"
"      - activation=softmax, loss=sparse_categorical_crossentropy\n"
"      - batch_size=32–128\n"
"      - dynamically compute top_k as: \n"
"      - num_classes = len(np.unique(y_enc)) if isinstance(y_enc, (list, np.ndarray))
else 3\n"
"      - top_k = min(num_classes, 5)\n"
"      - metrics = ['accuracy',
tf.keras.metrics.SparseTopKCategoricalAccuracy(k=top_k, name=f'top_{top_k}_accuracy')] #
use sparse version for integer labels \n"
"      - at inference: take the top-`top_k` softmax probabilities for submission\n"
"    * multi-label-classification:\n"
"      - activation=sigmoid, loss=binary_crossentropy\n"
"      - batch_size=64–256, metrics=['accuracy', tf.keras.metrics.Precision(),
tf.keras.metrics.Recall(), tfa.metrics.F1Score(num_classes=n_classes)]\n"
"    * regression:\n"
"      - activation=linear, loss=mean_squared_error\n"

```

```

        "    - batch_size=32-256,
metrics=[tf.keras.metrics.RootMeanSquaredError(name='rmse'),
tf.keras.metrics.MeanAbsoluteError(name='mae')]\n"
    "    * **time-series forecasting:**\n"
    "    - use chronological split\n"
    "    - **model:** stack LSTM layers (their internal activations are tanh + sigmoid
gates), **then** any Dense head\n"
    "    - activation=linear, loss=mean_squared_error\n"
    "    - epochs=10-50,
metrics=[tf.keras.metrics.RootMeanSquaredError(name='rmse')]\n"
    ")\n"
    "7. **Compile the model with the Adam optimizer and the chosen loss and metrics\n"
    "8. **Callbacks & Training**:\n"
    "    start_time = time.time() # capture before fit\n"
    "    if y_val is not None:\n"
    "        history = model.fit(X_train_proc, y_train, validation_data=(X_val_proc, y_val),
epochs=100, callbacks=callbacks, verbose=2)\n"
    "    else:\n"
    "        history = model.fit(X_train_proc, y_train, validation_split=0.2, epochs=100,
callbacks=callbacks, verbose=2)\n"
    "    duration = time.time() - start_time # compute after fit\n"
    "9. **Evaluation & Logging**:\n"
    "    Don't use tensorflow_addons, it is no longer supported use more recent ways to
record metrics"
    "    Turn on the verbose and save the training and validation accuracy and log of the last
epoch in a json file (results.json). It will have the following keys: {training_accuracy,
training_loss, validation_accuracy and validation_loss}\n"
    "    with open('results.json', 'w') as f: json.dump(results, f)\n"
    "# Infer id_col & target_columns from submission_example header\n
if any(not col.replace('.', '').isdigit() for col in target_columns) or len(target_columns) >
1:\n
    competition_problem_subtype = "multi-label-classification"\n
\n
10. **Prediction & Submission**:\n
raw_preds = model.predict(X_test_proc)\n
if competition_problem_subtype == "multi-label-classification": final = (raw_preds >
0.5).astype(int)\n
    elif competition_problem_subtype in ["multiclass", "multiclass-classification"]: idxs =
raw_preds.argmax(axis=1); final = le.inverse_transform(idxs)\n
    elif competition_problem_subtype == "binary-classification":\n
        probs = raw_preds[:, 1] if raw_preds.ndim==2 and raw_preds.shape[1]==2 else
raw_preds.flatten()\n
        final = (probs > 0.5).astype(int)\n

```

```

        elif competition_problem_subtype in
["continuous-regression","quantile-regression","multi-output
regression","missing-value-imputation"]:\n\
        final = raw_preds\n\
        if np.all(final >= 0): final = np.expm1(np.clip(final, a_min=None, a_max=20))\n\
        else: final = raw_preds\n\
\n\
        # ensure 2D\n\
        if final.ndim == 1: final = final.reshape(-1,1)\n\
        submission = pd.DataFrame(final, columns=target_columns)\n\
        submission.insert(0, id_col, test_ids.reset_index(drop=True))\n\
        submission.to_csv('submission_result.csv', index=False)\n"

),
"parameters": {
    "type": "object",
    "additionalProperties": False,
    "properties": {
        "competition_problem_description": {
            "type": "string",
            "description": "Dense competition description giving the core goal."
        },
        "competition_problem_subtype": {
            "type": "string",
            "enum": [
                "binary-classification",
                "multiclass-classification",
                "multi-label-classification",
                "time-series-forecasting",
                "continuous-regression",
                "quantile-regression",
                "multi-output-regression",
                "ordinal-regression",
                "missing-value-imputation"
            ],
            "description": "Rely on this to choose splits, loss, activation, etc."
        },
        "dataset_metadata": {
            "type": "string",
            "description": "Full NLP explanation of the dataset, the columns that need to be
predicted and the training files provided"
        },
        "data_profiles": {
            "type": "array",

```



```

"description": "One entry per file after compaction",
"items": {
  "type": "object",
  "required": ["file_name", "shape", "targets"],
  "properties": {
    "file_name": { "type": "string" },

    "shape": {
      "type": "object",
      "required": ["rows", "cols"],
      "properties": {
        "rows": { "type": "integer" },
        "cols": { "type": "integer" }
      },
      "additionalProperties": False
    },

    "targets": {
      "type": "array",
      "items": { "type": "string" }
    }
  },
  "additionalProperties": False
},
"files_preprocessing_instructions": {
  "type": "string",
  "description": "Instructions for how to preprocess the raw files."
},
"submission_example": {
  "type": "array",
  "minItems": 1,
  "items": {
    "type": "object",
    "properties": {
      "column_name": { "type": ["string", "number"] },
      "value": { "type": ["number", "string", "boolean", "null"] }
    },
    "required": ["column_name", "value"],
    "additionalProperties": False
  }
},
"files_list": {
  "type": "array",

```

```

        "items": {"type": "string"},
        "description": " list of all files included in the competition, decide whether there are
testing files and whether you need to split the training dataset"
    },
    "examples": {
        "type": "array",
        "description": "Retrieved preprocessing and code snippets from solutions of top
similar competitions, rely on them ",
        "items": {
            "type": "object",
            "additionalProperties": False,
            "properties": {
                "preprocessing_steps": {
                    "type": "array",
                    "items": {"type": "string"}
                },
                "model_layers_code": {"type": "string"}
            },
            "required": ["preprocessing_steps", "model_layers_code"]
        },
        "notebook_code": {
            "type": "string",
            "description": "****The complete runnable Python notebook code wrapped in
<Code>...</Code>."
        }
    },
    "required": [
        "competition_problem_description",
        "competition_problem_subtype",
        "dataset_metadata",
        "data_profiles",
        "files_preprocessing_instructions",
        "submission_example",
        "files_list",
        "examples",
        "notebook_code"
    ],
    "strict": True
}
]

```

Extracting the Keras model layer and fit block

```
extract_tools = [
    {
        "name": "extract_model_block",
        "type": "function",
        "description": (
            "Given the full notebook in `original_code`, return JSON with exactly two keys:\n"
            "- `model_block`: every line from the first `model =` up to the line **before** the\n"
            "`model.fit(...)` call;\n"
            "- `fit_call`: the **entire** `model.fit(...)` invocation (including all its arguments) as one\n"
            "string.\n"
            "Do not change anything—just extract those snippets."
        ),
        "parameters": {
            "type": "object",
            "properties": {
                "original_code": { "type": "string" },
                "model_block": { "type": "string" }
            },
            "required": ["original_code", "model_block"],
            "additionalProperties": False
        }
    }
]
```

#Building a Keras Tuner model using the previously extracted, working Keras model

```
tuner_tools = [  
    #Tuner  
    {  
        "name": "generate_tuner_schema",  
        "type": "function",  
        "description": (  
            """Return ONE JSON field `tuner_code` containing runnable Python wrapped in  
<Code>...</Code>."""  
            "1. Choose profile → best tag overlap in `hyperparameter_bank` → `chosen`.\n\n"  
            "• **Preserve** the definitions and usages of:\n"  
            "  ``python\n"  
            "    early_stopping = EarlyStopping(monitor='val_loss', patience=10,  
restore_best_weights=True)\n"  
            "    checkpoint    = ModelCheckpoint('best_model.h5', monitor='val_loss',  
save_best_only=True)\n"  
            "  ``\n"  
            " and plug these same variables into both `tuner.search(...)` and the final  
`model.fit(...)`.\n\n"  
            "# preserve input dim\n"  
            "n_features = X_train_proc.shape[1]\n\n"  
            "2. HyperModel\n``python\n"  
            "import keras_tuner as kt\n"  
            "from tensorflow.keras.layers import Input, Dense, Dropout\n"  
            "from tensorflow.keras.models import Model\n\n"  
            "class MyHyperModel(kt.HyperModel):\n"  
            "    def build(self, hp):\n"  
            "        layers = hp.Int('layers', **chosen['params']['layers'])\n"  
            "        units = hp.Int('units', **chosen['params']['units'])\n"  
            "        act = hp.Choice('activation', chosen['params']['activation']['values'])\n"  
            "        drop = hp.Float('dropout', **chosen['params']['dropout'])\n"  
            "        opt = hp.Choice('optimizer', chosen['params']['optimizer']['values'])\n"  
            "        lr = hp.Float('learning_rate', **chosen['params']['learning_rate'],  
sampling='log')\n\n"  
            "        inputs = Input(shape=(n_features,))\n"  
            "        x = inputs\n"  
            "        if 'lstm_units' in chosen['params']:\n"  
            "            dense_units = hp.Int('dense_units', **chosen['params']['dense_units'])\n"  
            "            # time-series LSTM branch\n"  
            "            for i in range(layers):\n"  
            "                return_seq = (i < layers - 1)\n"  
            "                x = LSTM(lstm_units, return_sequences=False)(x)
```

```

"        x = Dense(dense_units, activation='relu')(x)\n"
"        x = Dropout(drop)(x)\n"
"    else:\n"
"        for _ in range(layers):\n"
"            x = Dense(units, activation=act)(x)\n"
"            x = Dropout(drop)(x)\n"
"    x = output_layer_original(x) # keep orig head\n"
"    model = Model(inputs, x)\n"
"    model.compile(optimizer=opt, loss=original_loss, metrics=original_metrics)\n"
"    # stash for use in tuner.search\n"
"    return model\n"
""" \n"

```

"""No extra `hp.` calls."""\n\n

3. Replace `model.fit` with Bayesian only, **using the provided batch_size and epochs from the **hyperparameter_bank:\n"

```

"tuner = kt.BayesianOptimization(\n"
"    MyHyperModel(),\n"
"    objective='val_loss',\n"
"    max_trials=10,\n"
"    executions_per_trial=1,\n"
"    seed=42,\n"
"    overwrite=False,\n"
"    project_name='bayesian_tuner'\n"
")\n\n"
"if y_val is not None:\n"
"    tuner.search(\n"
"        X_train_proc, y_train,\n"
"        validation_data=(X_val_proc, y_val),\n"
"        batch_size=bs, epochs=ep,\n"
"        callbacks=[early_stopping, checkpoint]\n"
"    )\n"
"else:\n"
"    tuner.search(\n"
"        X_train_proc, y_train,\n"
"        validation_split=0.2,\n"
"        batch_size=bs, epochs=ep,\n"
"        callbacks=[early_stopping, checkpoint]\n"
"    )\n\n"
"model = tuner.hypermodel.build(\n"
"    tuner.get_best_hyperparameters(1)[0]\n"
")\n"
"""

```

4. Retrain `model` with the original callbacks and data, guarding against `None`:\n"

```

"""python\n"

```

```

    "if y_val is not None:\n"
    "    history = model.fit(\n"
    "        X_train_proc, y_train,\n"
    "        validation_data=(X_val_proc, y_val),\n"
    "        epochs=100, batch_size=bs,\n"
    "        callbacks=[early_stopping, checkpoint],\n"
    "        verbose=2\n"
    "    )\n"
    "else:\n"
    "    history = model.fit(\n"
    "        X_train_proc, y_train,\n"
    "        validation_split=0.2,\n"
    "        epochs=100, batch_size=bs,\n"
    "        callbacks=[early_stopping, checkpoint],\n"
    "        verbose=2\n"
    "    )\n"
    """\n"
),
"parameters": {
    "type": "object",
    "additionalProperties": False,
    "properties": {
        "competition_problem_description": {
            "type": "string",
            "description": "Full text description of the task."
        },
        "competition_problem_subtype": {
            "type": "string",
            "enum": [
                "binary-classification",
                "multiclass-classification",
                "multi-label-classification",
                "time-series-forecasting",
                "continuous-regression",
                "quantile-regression",
                "multi-output-regression",
                "ordinal-regression",
                "missing-value-imputation"
            ]
        },
        "model_block": {
            "type": "string",
            "description": "The code from the Keras version of `model =` through just before  

`model.fit`, plus any batch_size/epochs definitions."
        }
    }
}

```

```

    },
    "hyperparameter_bank": {
      "type": "object",
      "description": "One selected profile from HYPERPARAMETER_BANK for this
task.",
      "additionalProperties": False,
      "properties": {
        "params": {
          "type": "object",
          "additionalProperties": False,
          "properties": {
            "layers": { "type": "integer", "minimum": 1, "maximum": 8, "multipleOf": 1 },
            "units": { "type": "integer", "minimum": 64, "maximum": 1024, "multipleOf":
64 },
            "activation": { "type": "string", "enum": ["relu", "tanh"] },
            "dropout": { "type": "number", "minimum": 0.0, "maximum": 0.5 },
            "optimizer": { "type": "string", "enum": ["adam"] },
            "learning_rate": { "type": "number", "minimum": 1e-5, "maximum": 1e-2 },
            "batch_size": { "type": "integer", "enum": [32,64,128,256,512,1024] },
            "epochs": { "type": "integer", "minimum": 10, "maximum": 200, "multipleOf":
10 },
            "dense_units": { "type": "integer", "minimum": 16, "maximum": 512,
"multipleOf": 16 }
          },
        },
      },
      "required": ["dropout","optimizer","learning_rate","batch_size","epochs",
"activation","layers","units"]
    },
    "tuner_choice": {
      "type": "string",
      "enum": ["gridsearch","bayesian","hyperband"],
      "description": "Which Keras Tuner class to use."
    },
    "tuner_code": {
      "type": "string",
      "description": "****The complete runnable Python notebook code wrapped in
<Code>...</Code> saved into the `tuner_code` JSON field."
    },
  },
  "required": [
    "competition_problem_description",
    "competition_problem_subtype",
    "model_block",

```

```
        "hyperparameter_bank",  
        "tuner_choice",  
        "tuner_code"  
    ]  
}  
}
```


Merge the existing Keras outline with the new Keras Tuner model layer code

```
merge =
{
"role": "user",
  "content": (
    "Here is my full notebook:\n\n"
    "```python\n"
    f"{original_code}\n```\n\n"
    "And here is the new Keras-Tuner snippet (build, compile, search, retrain):\n\n"
    "```python\n"
    f"{tuner_snippet}\n```\n\n"
    "Please replace only the existing model-definition block—that is, every line \n"
    "from the first `model =` up to (but not including) the first `model.fit` call—with this\n"
    "Keras-Tuner snippet. \n"
    "***Keep** any variables it relies on (`n_features`, `n_classes`, `output_layer_original`,\n"
    "etc.) so it drops in cleanly, \n"
    "and do not touch imports, data loading, preprocessing, callbacks, logging, or the\n"
    "submission code. Return the full notebook text with only that block swapped out."
  )
}
```